Elective Course

Course Code: CS4103

Autumn 2025-26

Lecture #14

Artificial Intelligence for Data Science

Week-4: PROBLEM SOLVING BY SEARCH

Adversarial Search Problem--- Games [Part-II]

Course Instructor:

Dr. Monidipa Das

Assistant Professor

Department of Computational and Data Sciences

Indian Institute of Science Education and Research Kolkata, India 741246

Minimax Example (revisited)

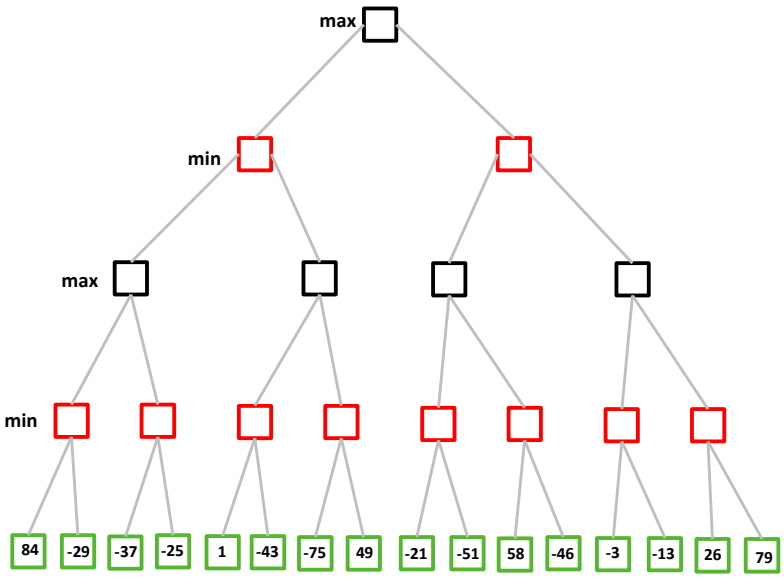

Minimax Example (revisited)

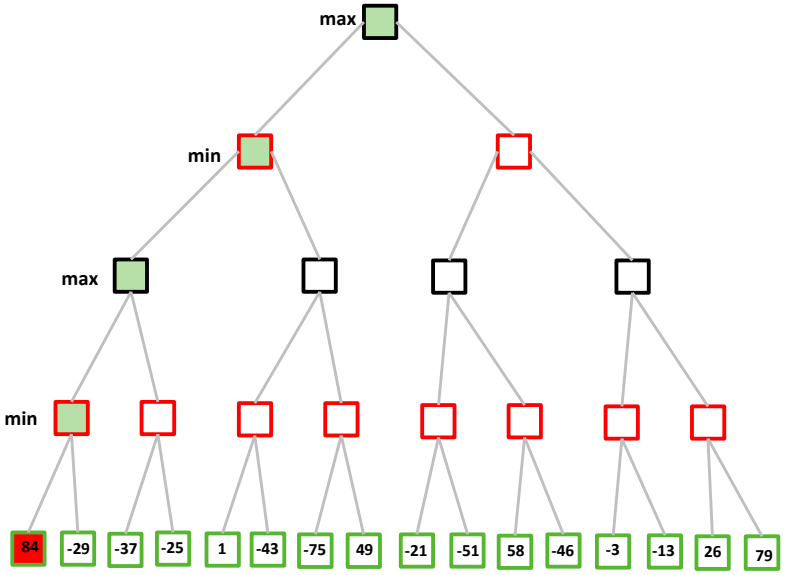

Minimax Example (revisited)

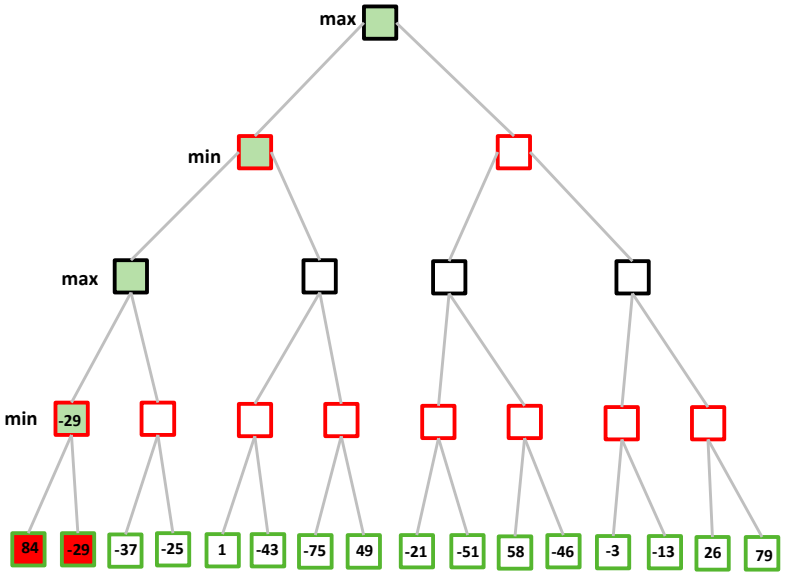

Minimax Example (revisited)

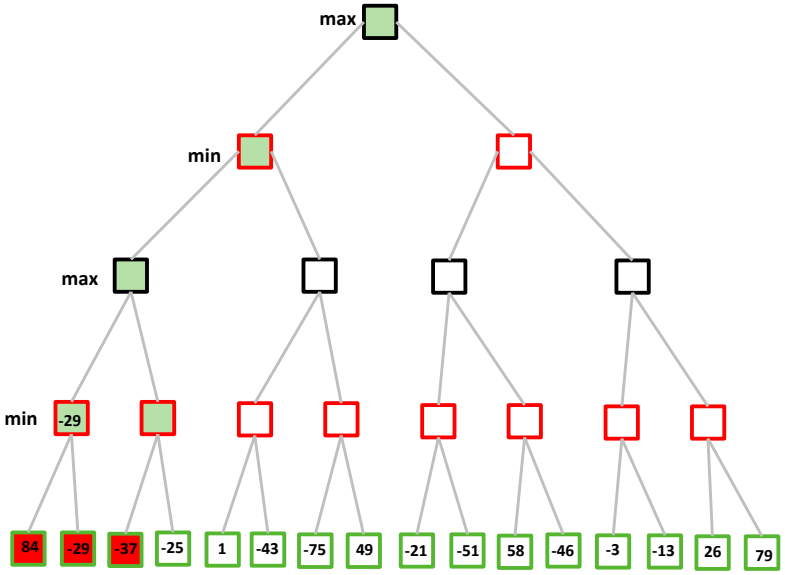

Minimax Example (revisited)

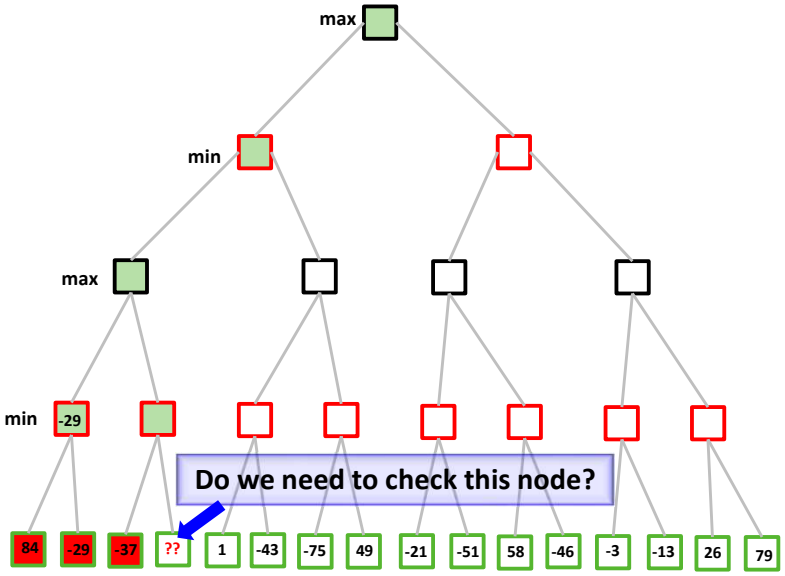

Minimax Example (revisited)

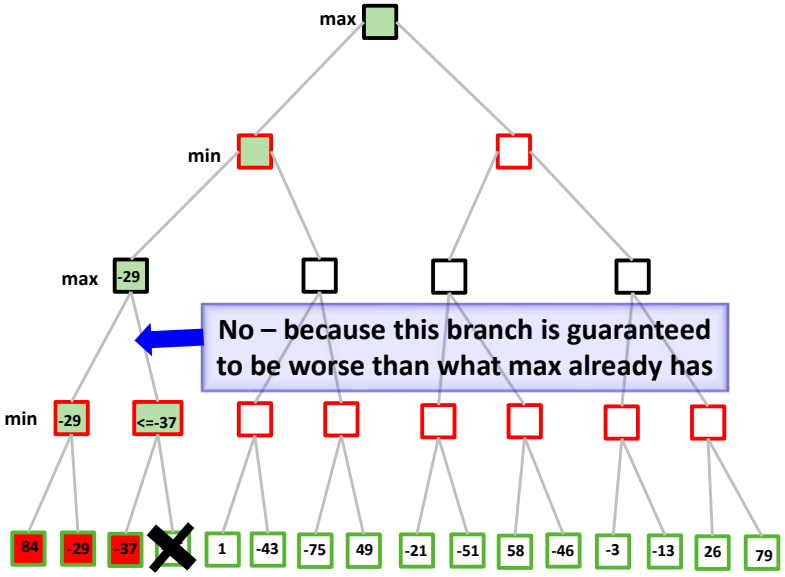

Alpha-Beta Procedure

- The alpha-beta procedure can speed up a depth-first minimax search.
- **Alpha:** a lower bound on the value that a max node may ultimately be assigned

$$v \geq \alpha$$

- **Beta:** an upper bound on the value that a min node may ultimately be assigned

$$v \leq \beta$$

Alpha-Beta Pruning

9

- Traverse the search tree in depth-first order
- At each MAX node n, $alpha(n)$ = maximum value found so far
- At each MIN node n, $beta(n)$ = minimum value found so far
  - **Note:** The alpha values start at -infinity and only increase, while beta values start at +infinity and only decrease.
- **Beta cutoff:** Given a MAX node n, cut off the search below n (i.e., don't generate or examine any more of n's children) if $alpha(n)$ >= $beta(i)$ for some MIN node ancestor i of n.
- **Alpha cutoff:** stop searching below MIN node n if $beta(n)$ <= $alpha(i)$ for some MAX node ancestor i of n.

Alpha Beta Example

10

- Do depth-first search until first leaf

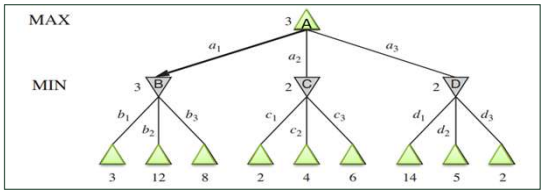

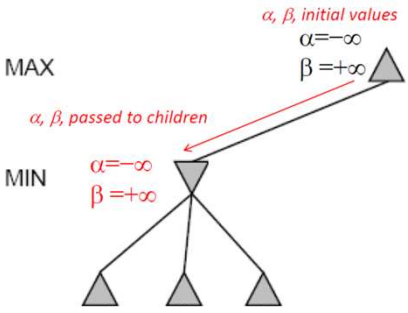

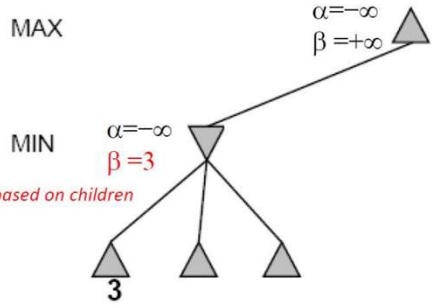

Alpha Beta Example

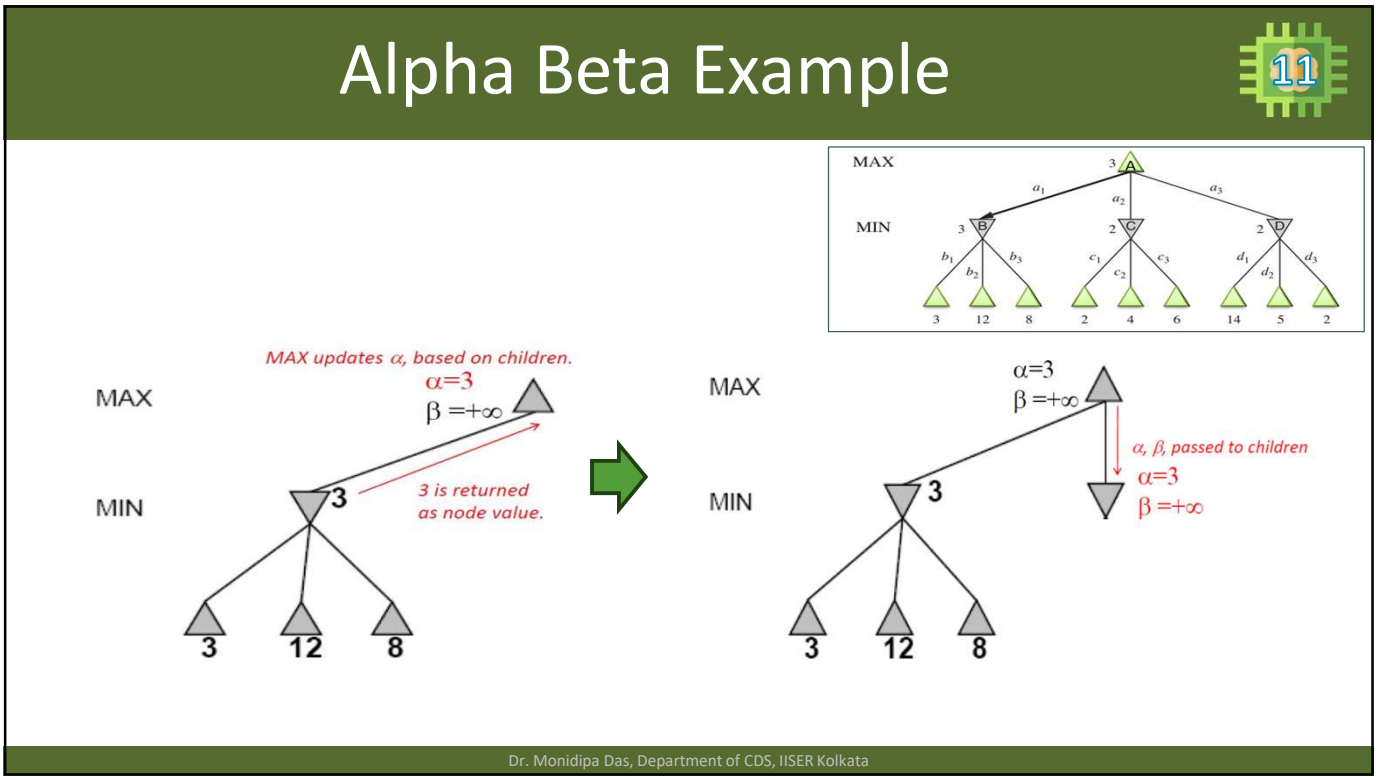

Alpha Beta Example

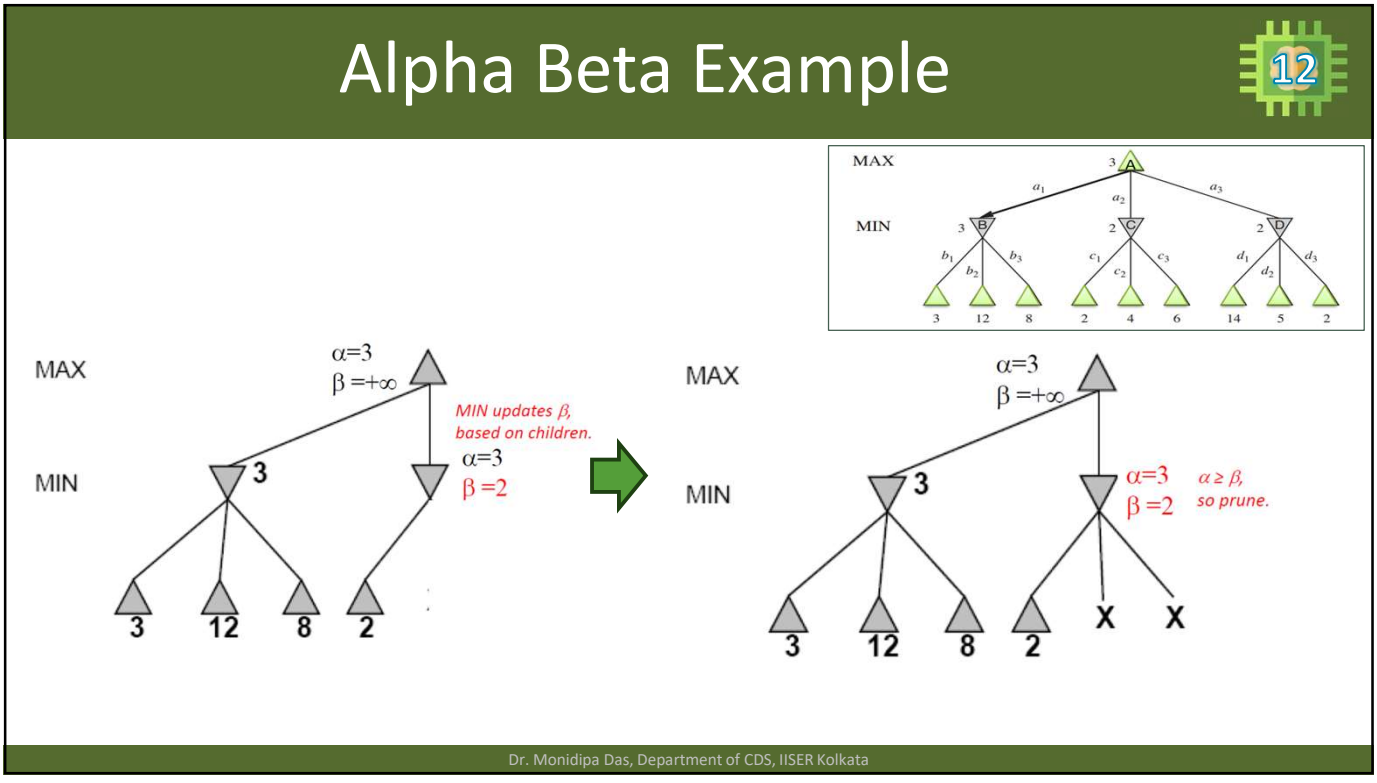

Alpha Beta Example

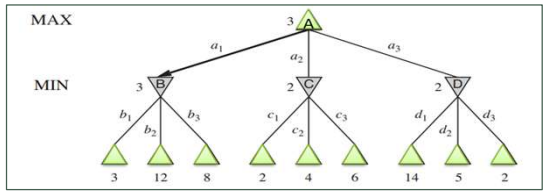

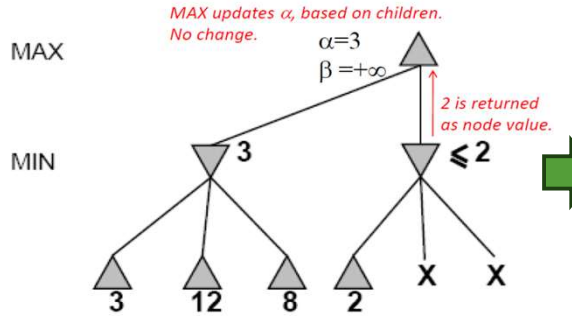

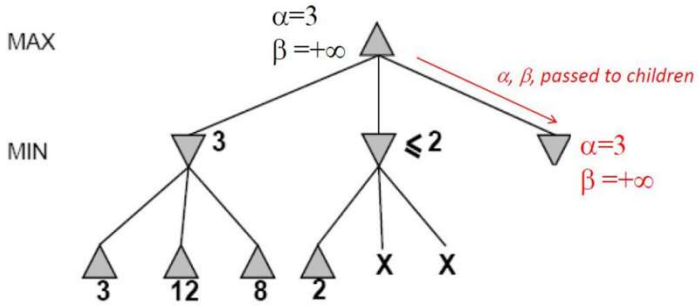

Alpha Beta Example

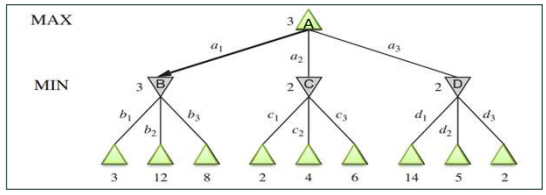

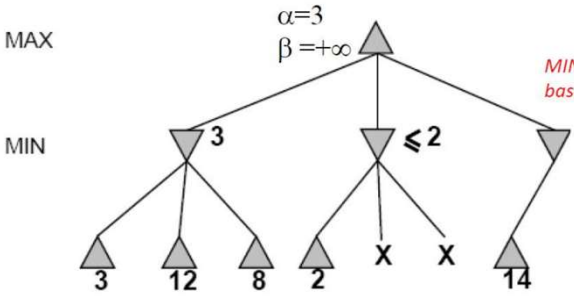

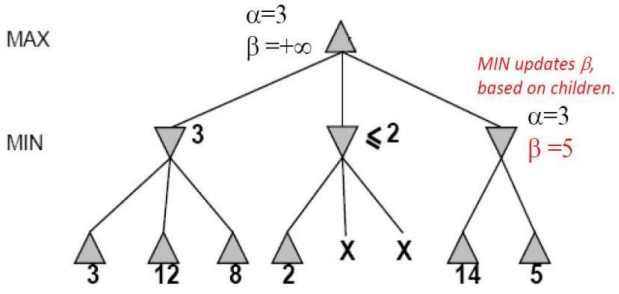

Alpha Beta Example

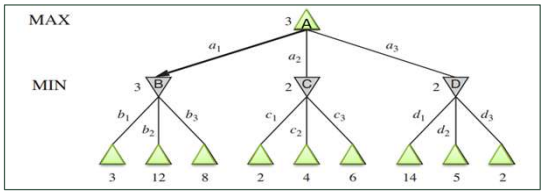

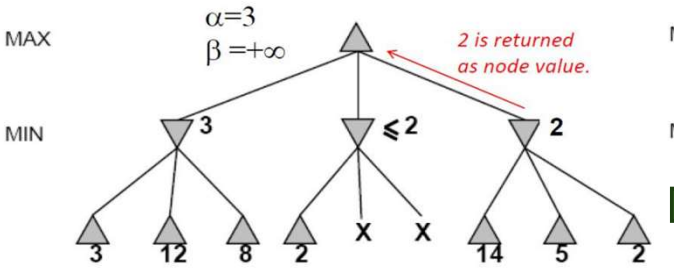

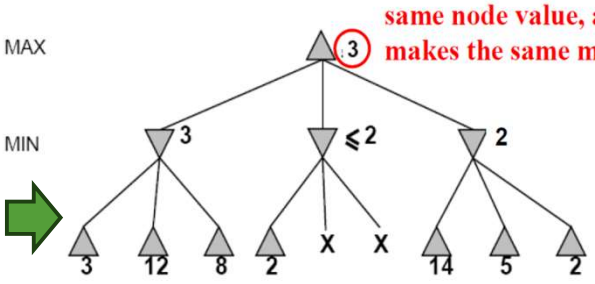

Max calculates the same node value, and makes the same move!

Alpha–Beta Search Algorithm

```
function ALPHA-BETA-SEARCH(state) returns an action
  v ← MAX-VALUE(state, −∞, +∞)
  return the action in ACTIONS(state) with value v

function MAX-VALUE(state, α, β) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
  v ← −∞
  for each a in ACTIONS(state) do
    v ← MAX(v, MIN-VALUE(RESULT(s,a), α, β))
    if v ≥ β then return v
    α ← MAX(α, v)
  return v

function MIN-VALUE(state, α, β) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
  v ← +∞
  for each a in ACTIONS(state) do
    v ← MIN(v, MAX-VALUE(RESULT(s,a), α, β))
    if v ≤ α then return v
    β ← MIN(β, v)
  return v
```

// no change of β value within MAX-VALUE()

$\alpha \leftarrow MAX(\alpha, v)$
$if\ \alpha \geq \beta\ then\ return\ v$

// no change of α value within MIN-VALUE()

$\beta \leftarrow MIN(\beta, v)$
$if\ \beta \leq \alpha\ then\ return\ v$

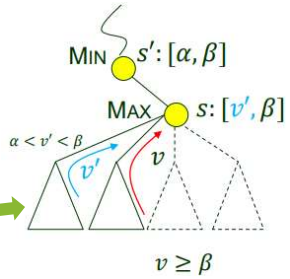

$v \geq \beta$

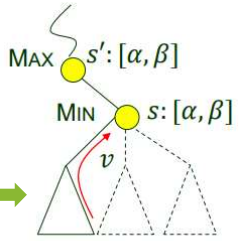

$v \leq \alpha$

Example

α - the best value for max along the path

β - the best value for min along the path

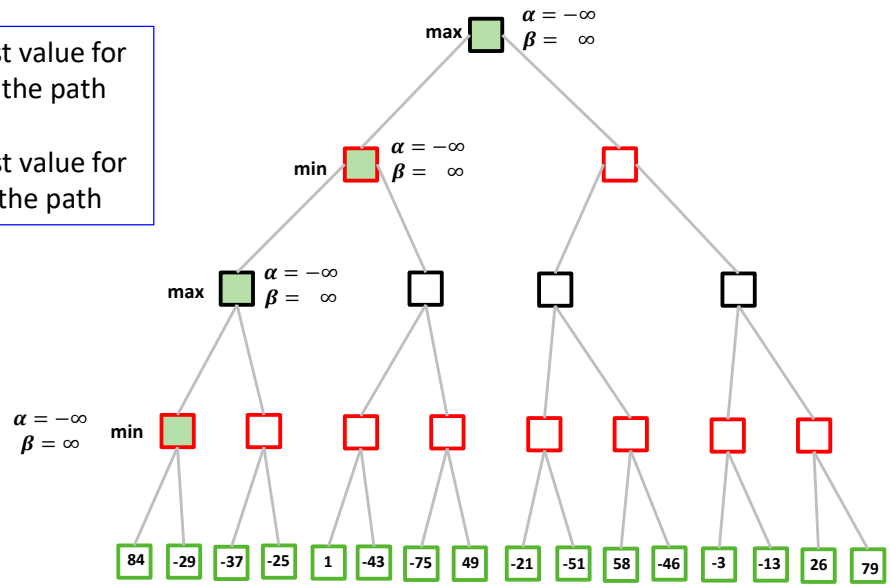

Example

α - the best value for max along the path

β - the best value for min along the path

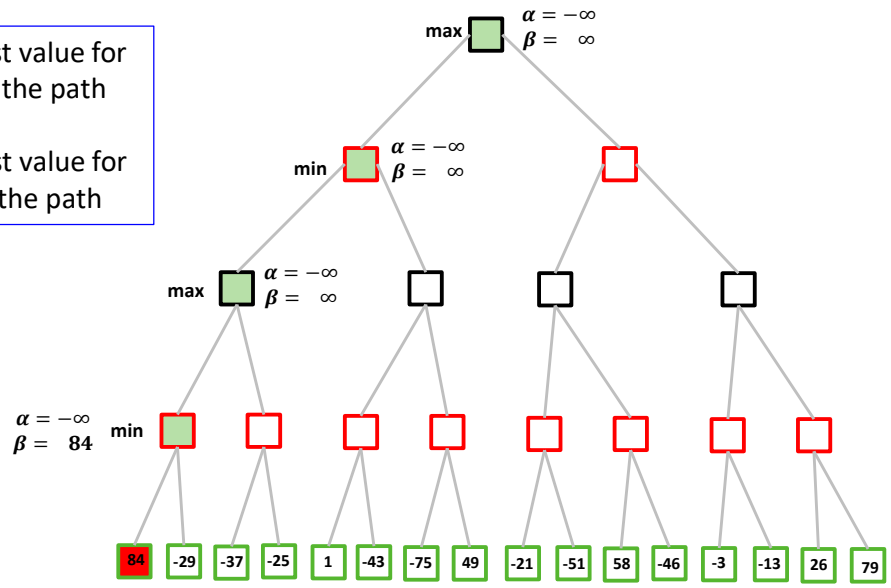

Example

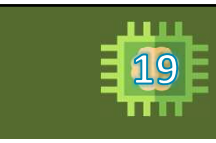

α - the best value for max along the path

β - the best value for min along the path

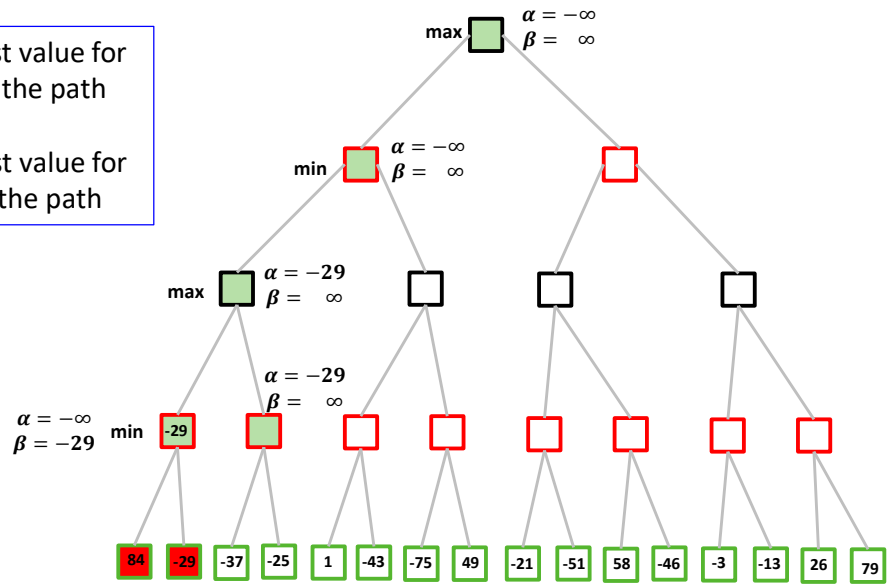

Example

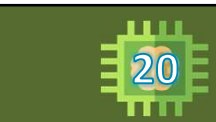

α - the best value for max along the path

β - the best value for min along the path

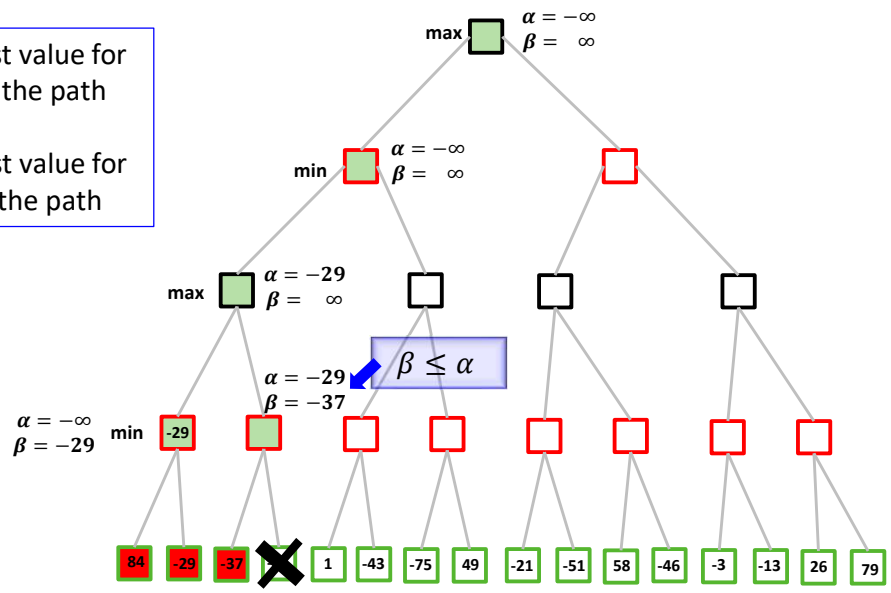

Example

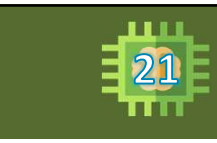

α - the best value for max along the path

β - the best value for min along the path

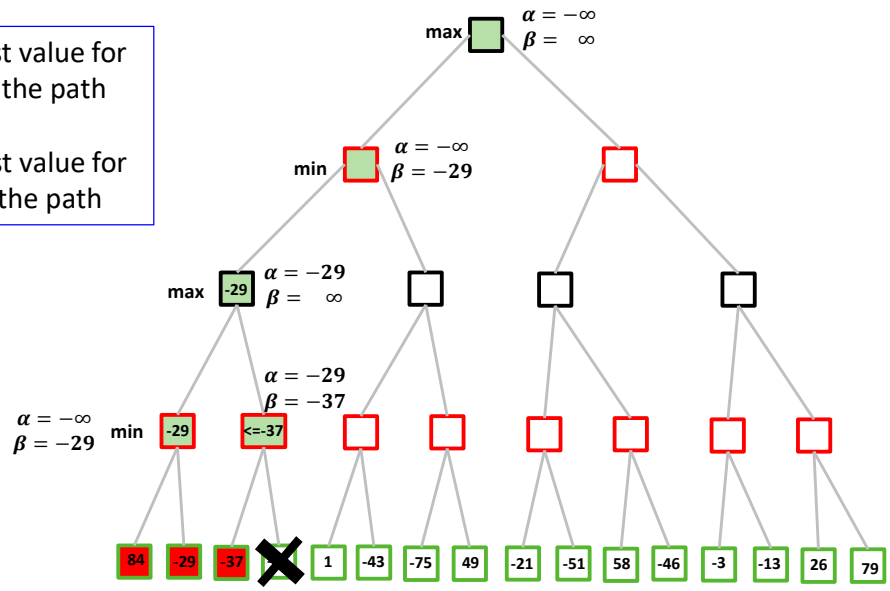

Example

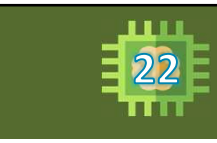

α - the best value for max along the path

β - the best value for min along the path

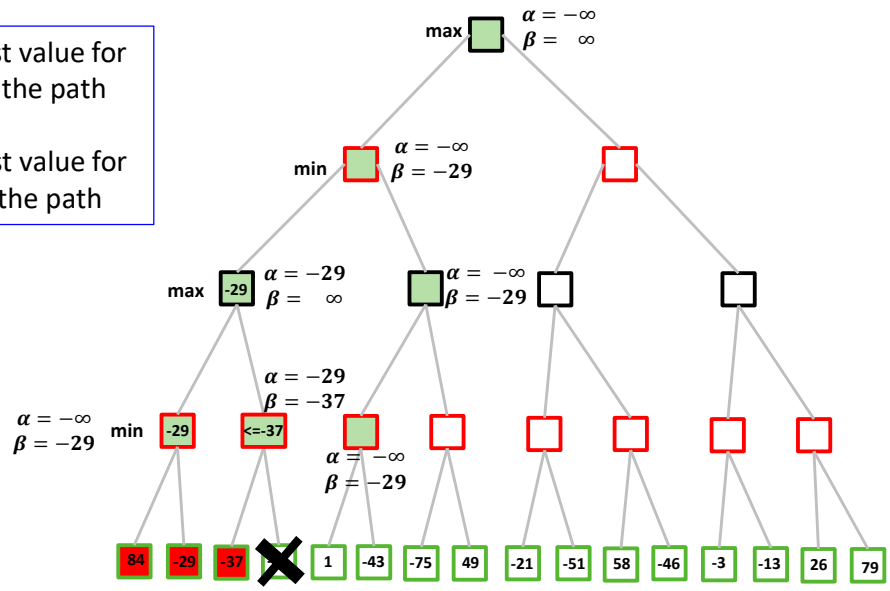

Example

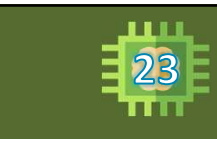

α - the best value for max along the path

β - the best value for min along the path

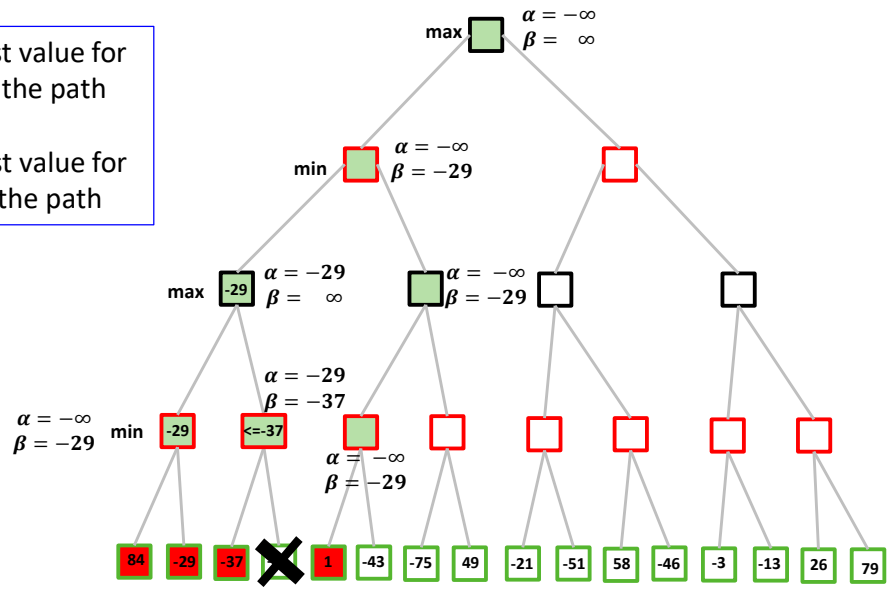

Example

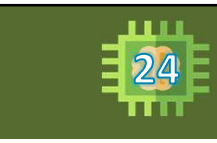

α - the best value for max along the path

β - the best value for min along the path

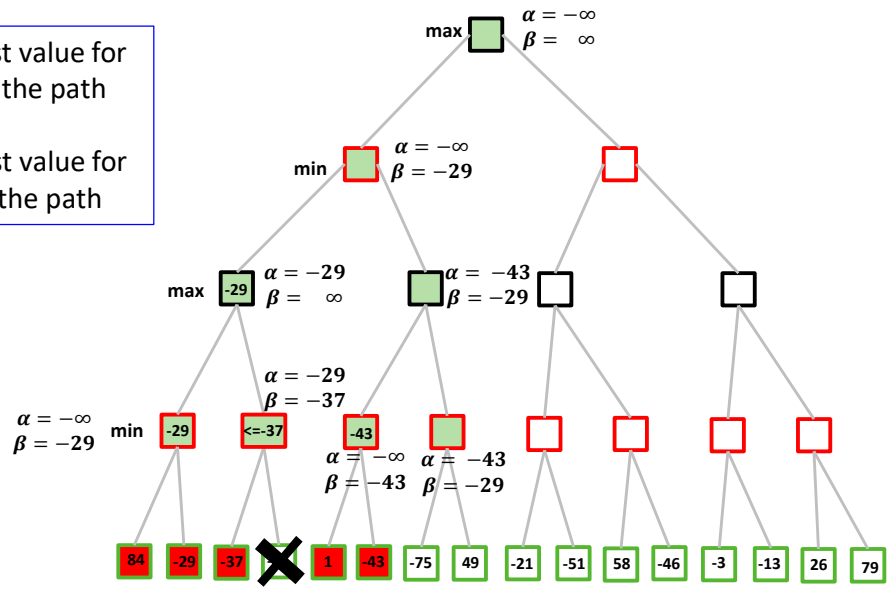

Example

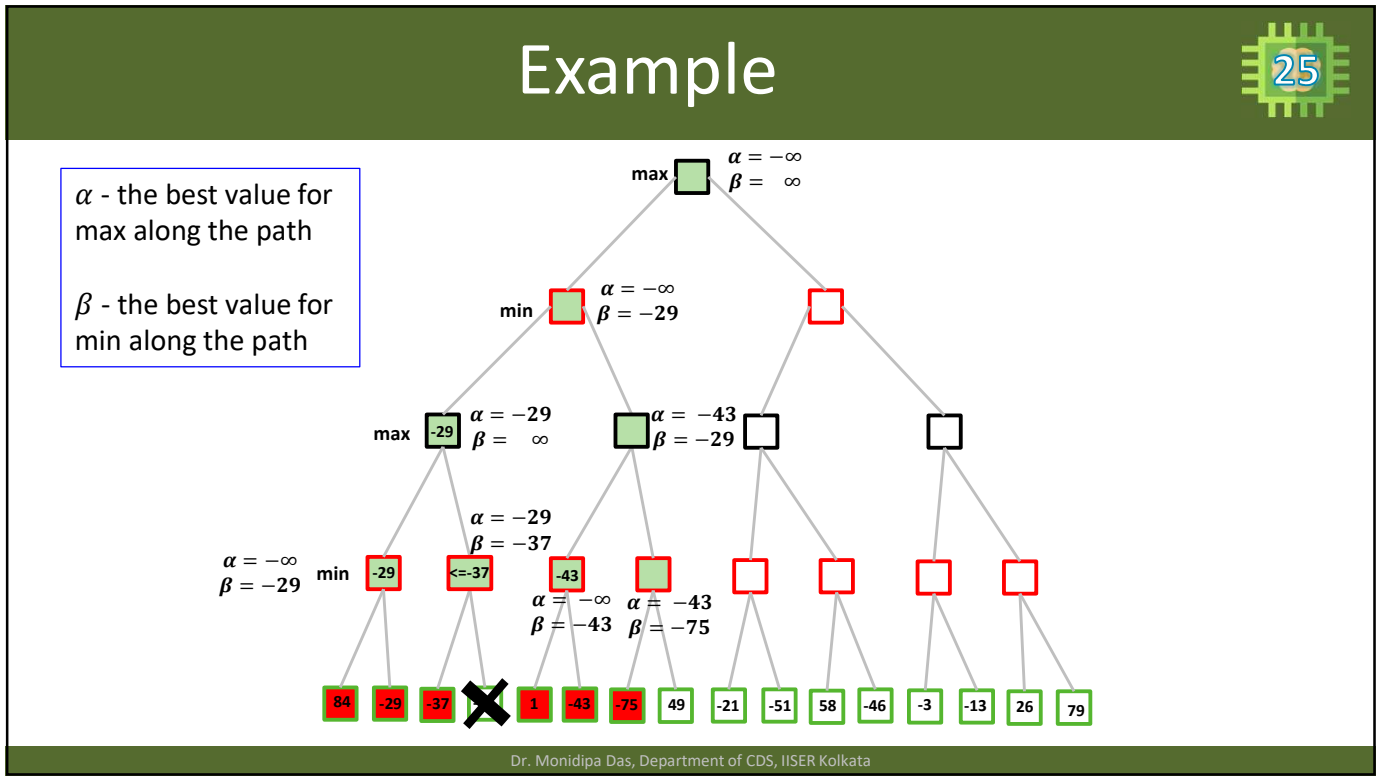

Example

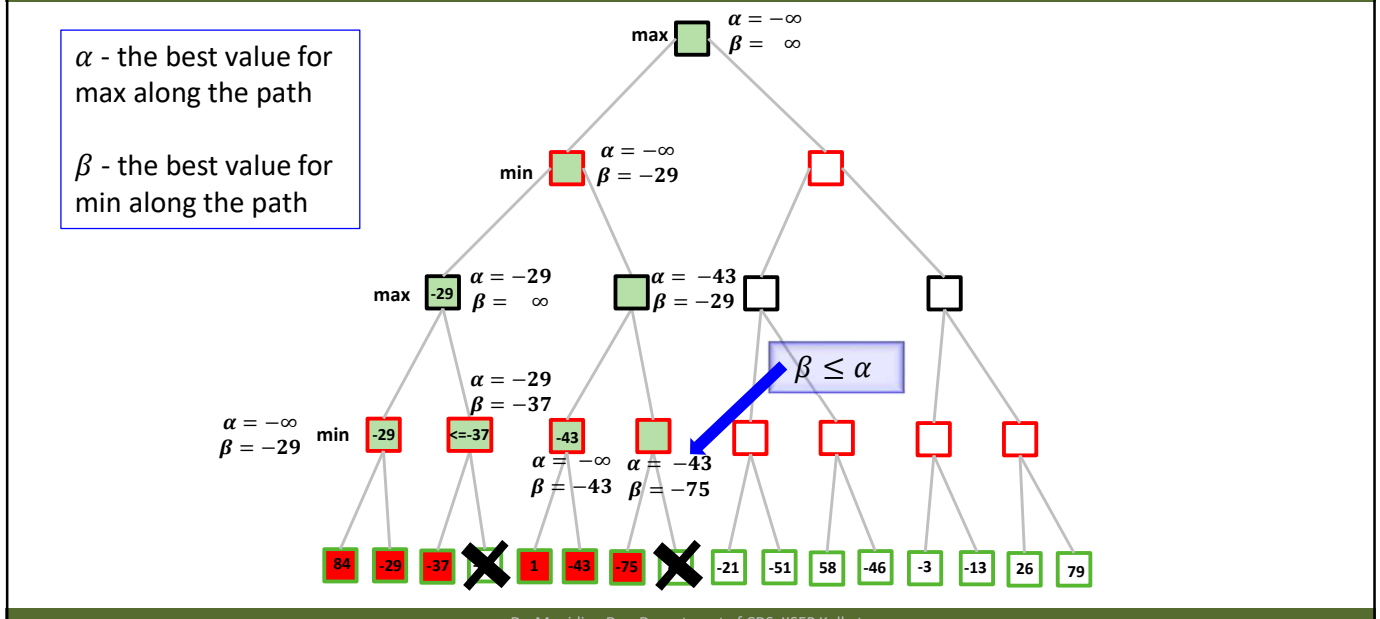

Example

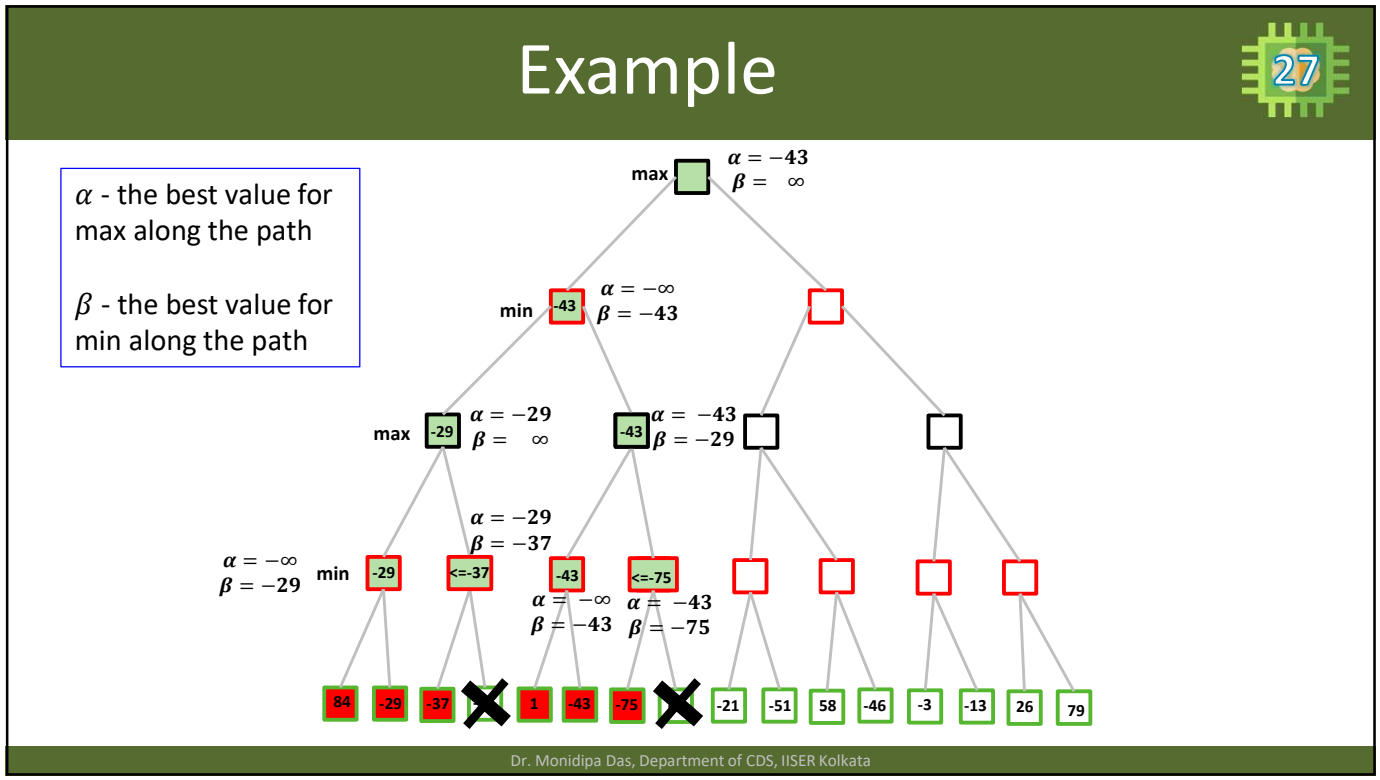

Example

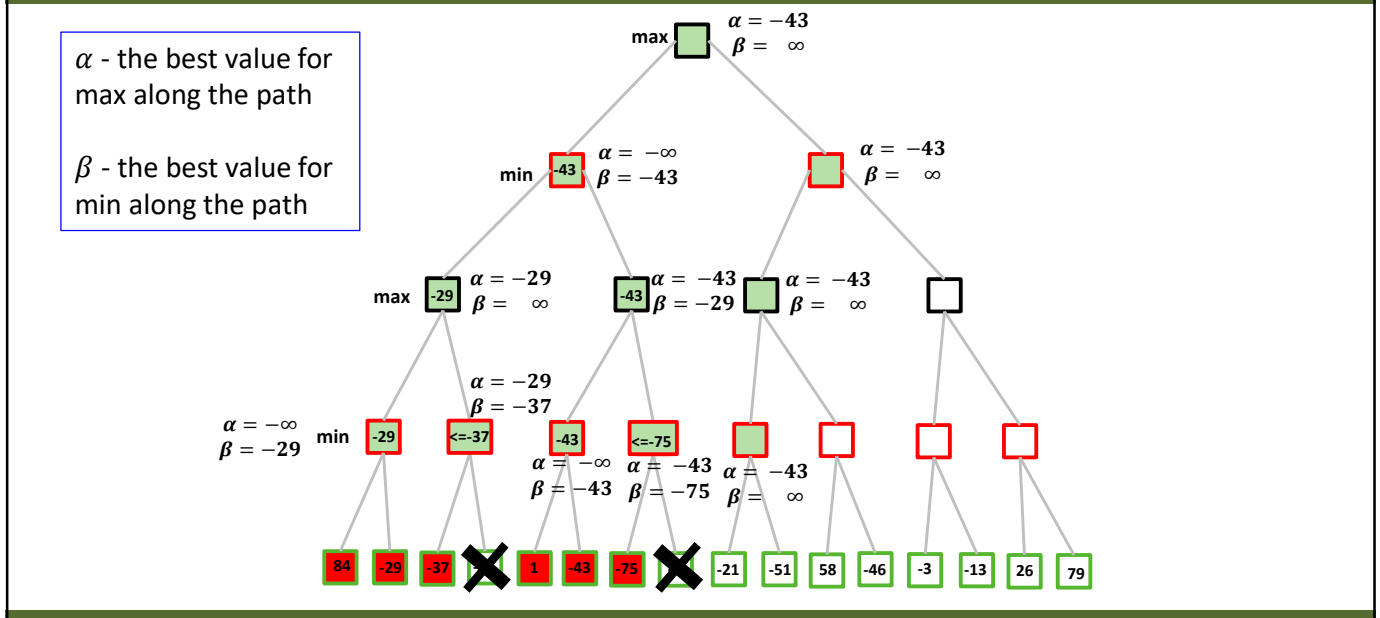

Example

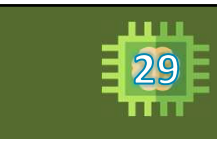

α - the best value for max along the path

β - the best value for min along the path

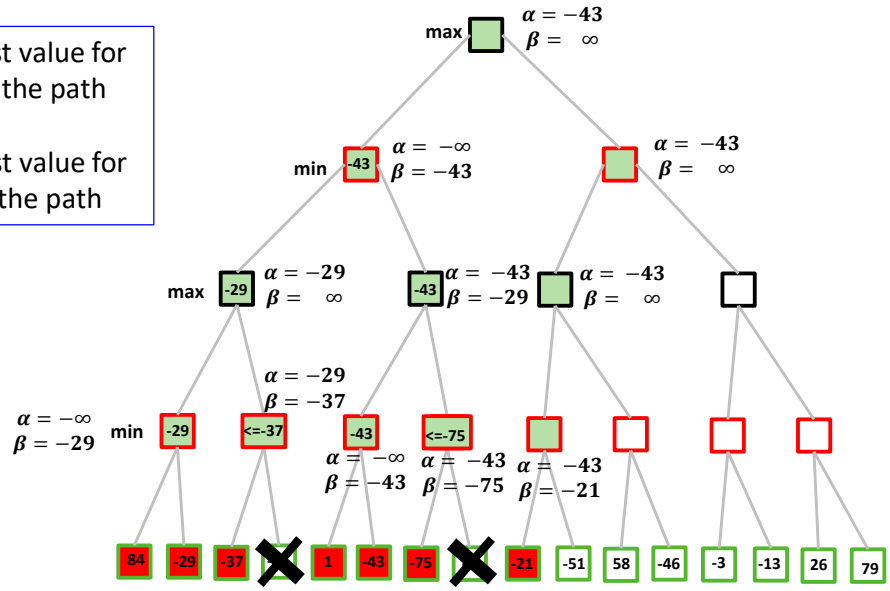

Example

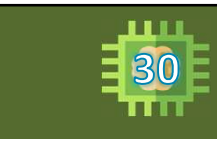

α - the best value for max along the path

β - the best value for min along the path

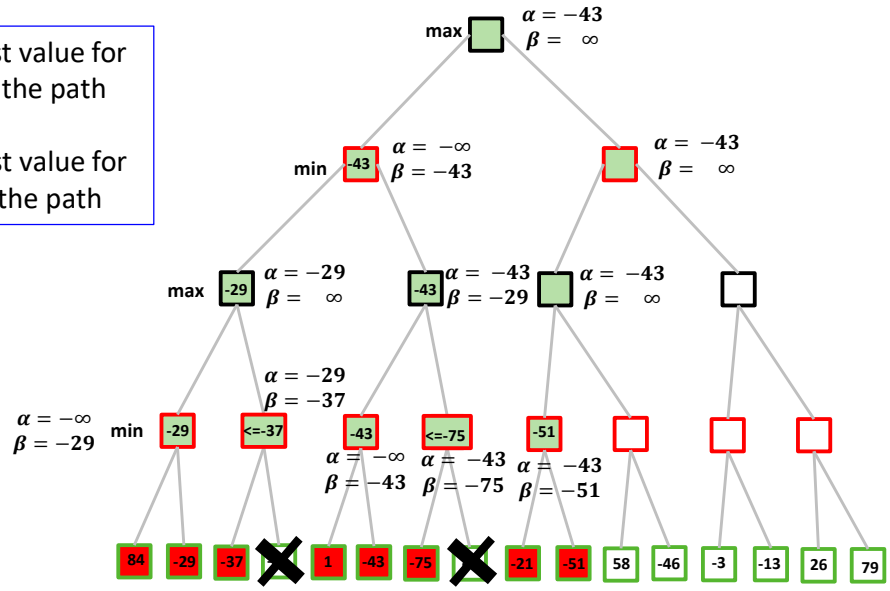

Example

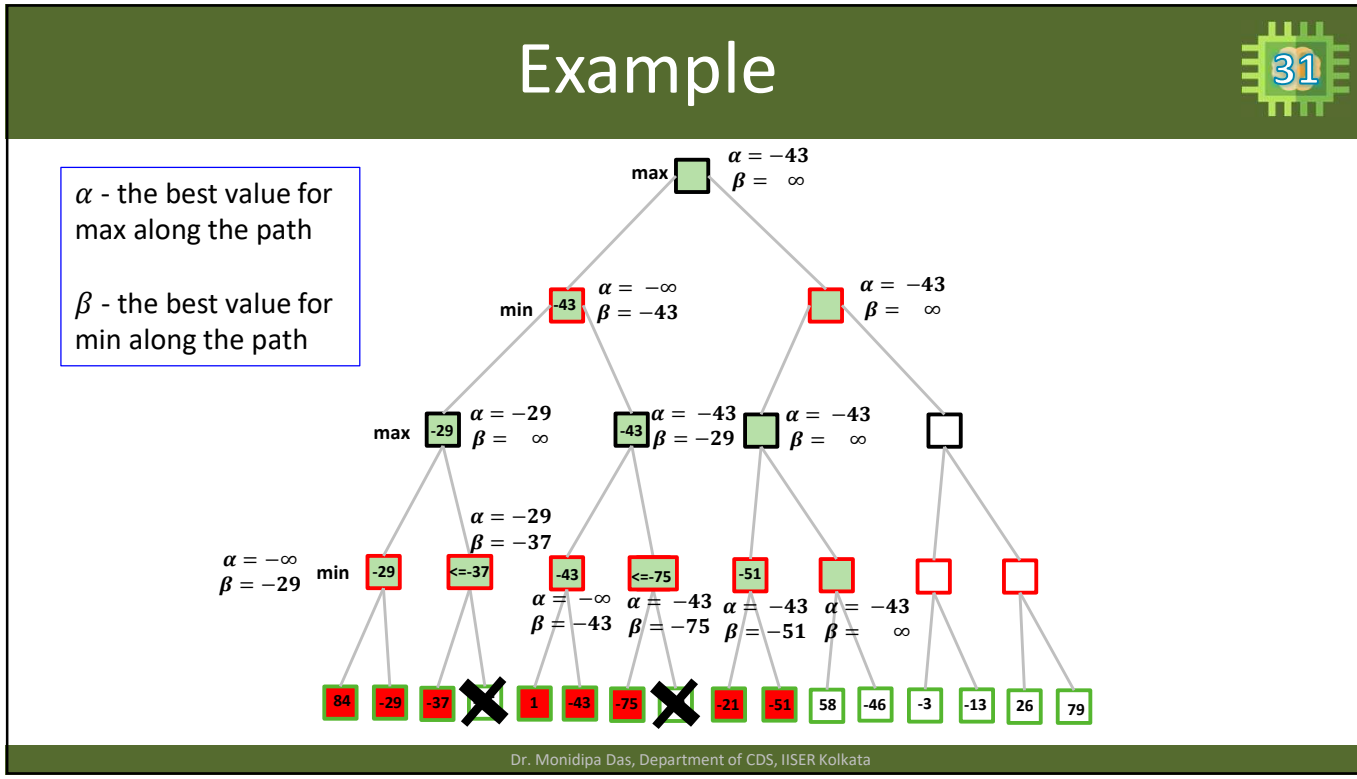

Example

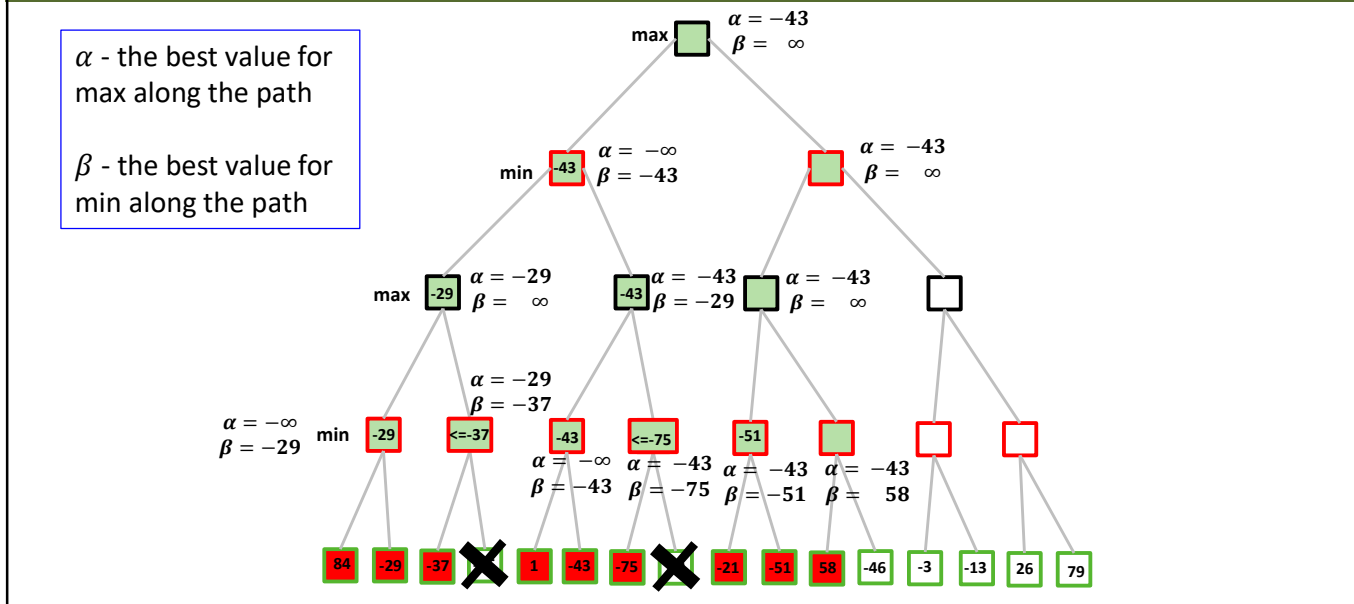

Example

33

α - the best value for max along the path

β - the best value for min along the path

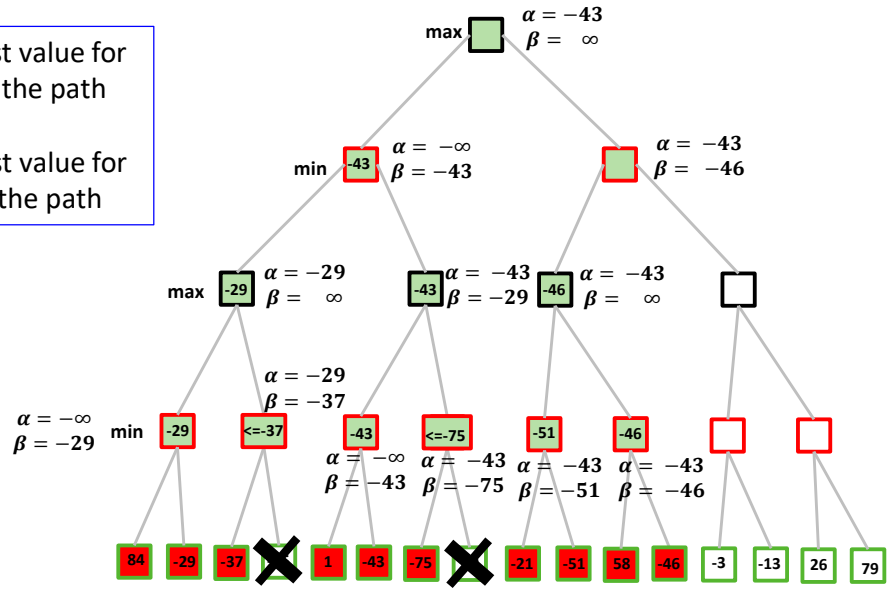

Example

34

α - the best value for max along the path

β - the best value for min along the path

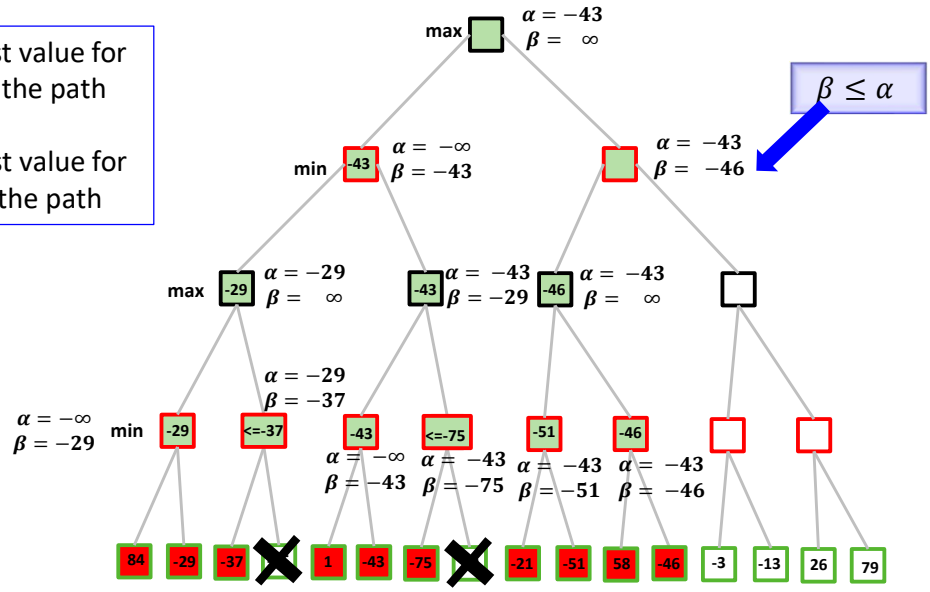

Example

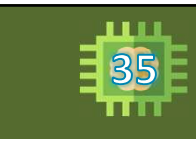

α - the best value for max along the path

β - the best value for min along the path

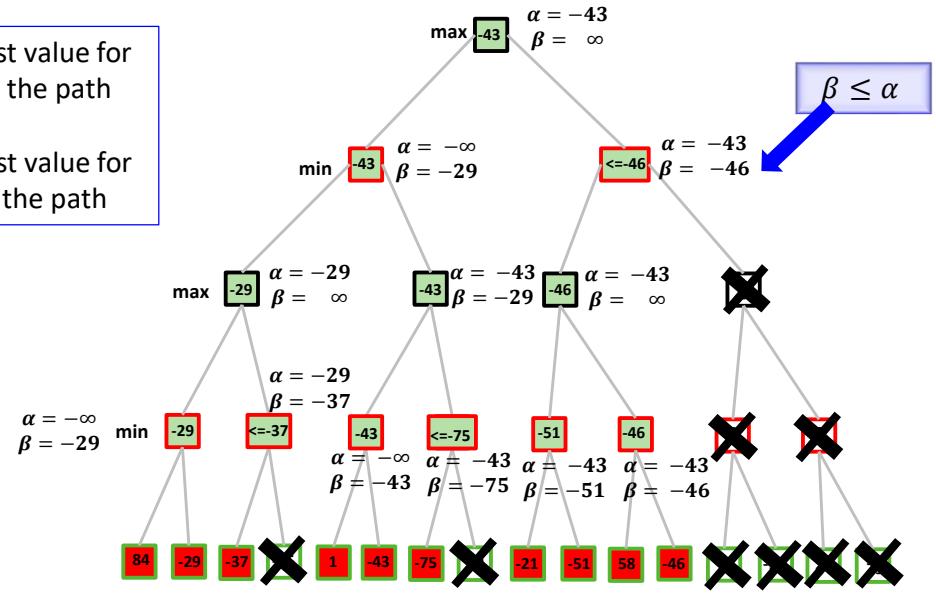

Example

α - the best value for max along the path

β - the best value for min along the path

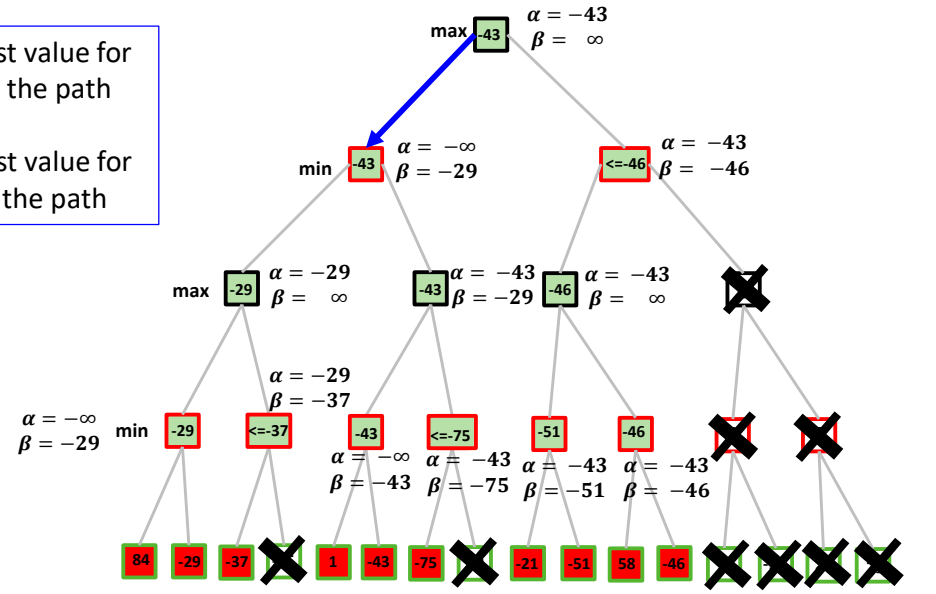

Effectiveness of alpha-beta

- Alpha-beta is guaranteed to compute the same value for the root node as computed by minimax, with less or equal computation
- **Worst case:** no pruning, examining b^m leaf nodes, where each node has b children and a m-ply search is performed
- **Best case:** examines only $O(b^{m/2})$ nodes.
  - Result is you can search twice as deep as minimax!
- **Best case** is when each player's best move is the first alternative generated

Move Ordering

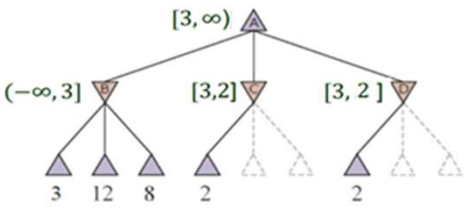

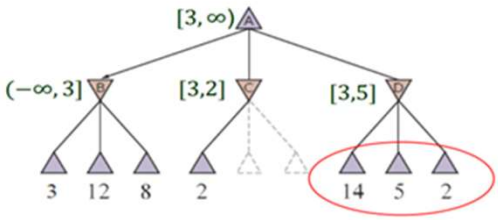

Successors 14 and 5 would've been pruned had 2 been generated first.

- Effectiveness of pruning is highly dependent on the order in which successors are generated.
- "Perfect ordering" has effective branching factor $\sqrt{b}$, which limits examination to only $O(b^{m/2})$ nodes compared to $O(b^m)$ for minimax.
- $O(b^{3m/4})$ nodes for random move ordering.

Good Enough?

- Chess (Minimax):
  - branching factor b≈35
  - game length m≈100
  - search space b m ≈ $35^{100} \approx 10^{154}$
  - Exact solution completely infeasible

- Chess (Alpha Beta Pruning):
  - branching factor b≈35
  - game length m≈100
  - search space $b^{m/2} \approx 35^{50} \approx 10^{77}$

Questions?